

Practice

- The gId is calculated from our configs, and therefore inactive threads do not interfere in the calculation of gId's. If you launch 2 blocks with 40 threads, even if you use 4 warps, the thread gIds will go from 0 to 79, for example.

Exercise

Now we will process our image of 768 x 448 pixels, divided into 14,24 blocks. For now, the image is read as grayscale.

Solution

```
1  #include "cuda_runtime.h"
2  #include "device_launch_parameters.h"
3
4  #include <stdio.h>
5  #include <stdlib.h>
6  #include <opencv2/opencv.hpp>
7
8  __host__ void checkCUDAError(const char* msg) {
9      cudaError_t error;
10     cudaDeviceSynchronize();
11     error = cudaGetLastError();
12     if (error != cudaSuccess) {
13         printf("ERROR %d: %s (%s)\n", error,
14             cudaGetErrorString(error), msg);
15     }
16
17     __global__ void complement(uchar* dev_a, uchar* dev_b) {
18         // locate my current block row
19         int threads_per_block = blockDim.x * blockDim.y;
20         int threads_per_row = threads_per_block * gridDim.x;
21         int row_offset = threads_per_row * blockIdx.y;
22
23         // locate my current block column
24         int block_offset = blockIdx.x * threads_per_block;
25         int threadId_inside = blockDim.x * threadIdx.y + threadIdx
26             .x;
27
28         int gId = row_offset + block_offset + threadId_inside;
29         dev_b[gId] = 255 - dev_a[gId];
30     }
31
32     using namespace cv;
33     int main() {
```

```
33
34     Mat img = imread("antenaParalelo.jpg", IMREAD_GRAYSCALE);
35
36     const int R = img.rows;
37     const int C = img.cols;
38
39     Mat imgResult(img.rows, img.cols, img.type());
40     uchar* host_a, * host_b, * dev_a, * dev_b, * pImg;
41     host_a = (uchar*)malloc(sizeof(uchar) * R * C);
42     host_b = (uchar*)malloc(sizeof(uchar) * R * C);
43     cudaMalloc((void**)&dev_a, sizeof(uchar) * R * C);
44     checkCUDAEError("Error at malloc dev_a");
45     cudaMalloc((void**)&dev_b, sizeof(uchar) * R * C);
46     checkCUDAEError("Error at malloc dev_b");
47
48     // matrix as vector
49     for (int i = 0; i < R; i++) {
50         pImg = img.ptr<uchar>(i); // points to a row each time
51         for (int j = 0; j < C; j++) {
52             host_a[i * C + j] = pImg[j];
53         }
54     }
55     cudaMemcpy(dev_a, host_a, sizeof(uchar) * R * C,
56               cudaMemcpyHostToDevice);
57
58     dim3 block(32, 32);
59     dim3 grid(C / 32, R / 32);
60
61     complement << < grid, block >> > (dev_a, dev_b);
62     checkCUDAEError("Error at kernel");
63
64     cudaMemcpy(host_b, dev_b, sizeof(uchar) * R * C,
65               cudaMemcpyDeviceToHost);
66
67     for (int i = 0; i < R; i++) {
68         pImg = imgResult.ptr<uchar>(i);
69         for (int j = 0; j < C; j++) {
70             pImg[j] = host_b[i * C + j];
71         }
72     }
73
74     imshow("Image", img);
75     imshow("Image Result", imgResult);
76     waitKey(0);
77
78     free(host_a);
79     free(host_b);
80     cudaFree(dev_a);
81     cudaFree(dev_b);
82
83     return 0;
```

82 }