

Practice

Exercise 1

- Function that solve a system of linear equations in the host.
- Function that solves a system of linear equations in the device through the launch of a kernel (1 block and 1 thread).

The kernel must receive all coefficients as a vector (of size 6).

A linear system with the form:

$$ax + by = c \quad dx + ey = f,$$

Can be solved by the formulas:

$$x = (ce - bf) / (ae - bd) \quad y = (af - cd) / (ae - bd)$$

```
1  #include "cuda_runtime.h"
2  #include "device_launch_parameters.h"
3
4  #include <stdio.h>
5  #include <stdlib.h>
6
7  __host__ void linearSolveCPU(float* n, float* x, float* y) {
8      *x = (n[2] * n[4] - n[1] * n[5]) / (n[0] * n[4] - n[1] * n
9          [3]);
10     *y = (n[0] * n[5] - n[2] * n[3]) / (n[0] * n[4] - n[1] * n
11        [3]);
12 }
13
14 __global__ void linearSolveGPU(float* n, float* x, float* y)
15 {
16     *x = (n[2] * n[4] - n[1] * n[5]) / (n[0] * n[4] - n[1] * n
17        [3]);
18     *y = (n[0] * n[5] - n[2] * n[3]) / (n[0] * n[4] - n[1] * n
19        [3]);
20 }
21
22 int main()
23 {
24     float* n_host = (float*)malloc(sizeof(float) * 6); // if
25     // malloc, you need to initialize all spaces one by one
26     float* x_host = (float*)malloc(sizeof(float));
27     float* y_host = (float*)malloc(sizeof(float));
28
29     float* x_gpu = (float*)malloc(sizeof(float));
30     float* y_gpu = (float*)malloc(sizeof(float));
```

```
27     float* n_device;
28     float* x_device;
29     float* y_device;
30
31     cudaMalloc((void**)&n_device, sizeof(float) * 6);
32     cudaMalloc((void**)&x_device, sizeof(float));
33     cudaMalloc((void**)&y_device, sizeof(float));
34
35     n_host[0] = 5;
36     n_host[1] = 1;
37     n_host[2] = 4;
38     n_host[3] = 2;
39     n_host[4] = -3;
40     n_host[5] = 5;
41
42     *x_host = 0;
43     *y_host = 0;
44     *x_gpu = 0;
45     *y_gpu = 0;
46
47     cudaMemcpy(n_device, n_host, sizeof(float) * 6,
48               cudaMemcpyHostToDevice);
49     cudaMemcpy(x_device, x_host, sizeof(float),
50               cudaMemcpyHostToDevice);
51     cudaMemcpy(y_device, y_host, sizeof(float),
52               cudaMemcpyHostToDevice);
53
54     linearSolveCPU(n_host, x_host, y_host);
55     printf("CPU result \n");
56     printf("x = %f y = %f \n", *x_host, *y_host);
57
58     linearSolveGPU <<< 1, 1 >>> (n_device, x_device, y_device)
59     ;
60     cudaMemcpy(x_gpu, x_device, sizeof(float),
61               cudaMemcpyDeviceToHost);
62     cudaMemcpy(y_gpu, y_device, sizeof(float),
63               cudaMemcpyDeviceToHost);
64     printf("GPU result \n");
65     printf("x = %f y = %f \n", *x_gpu, *y_gpu);
66
67     free(n_host);
68     free(x_host);
69     free(y_host);
70     free(x_gpu);
71     free(y_gpu);
72
73     cudaFree(n_device);
74     cudaFree(x_device);
75     cudaFree(y_device);
76
77     return 0;
```

```
72 }
```

Lab 03

Make a program in c/c++ in which you launch a kernel with one block and one thread. The kernel must solve a quadratic equation in the form:

$$ax^2 + bx + c = 0,$$

where its solutions are given by:

$$x1 = (-b + \sqrt{b^2 - 4ac}) / 2a \quad x2 = (-b - \sqrt{b^2 - 4ac}) / 2a$$

For the implementation, you must consider:

1. Ask the user for coefficients a, b and c.
2. The program must show the solutions for the equation or a message stating that the solution does NOT exist if the result is an imaginary number.

Tests

- a = 1, b = -5, c = 6 -> x1 = 2, x2 = 3
- a = 1, b = 1, c = 1 -> The solution does not exist

Solution

```
1  #include "cuda_runtime.h"
2  #include "device_launch_parameters.h"
3
4  #include <stdio.h>
5  #include <stdlib.h>
6  #include <math.h>
7
8  __global__ void solveGPU(double* dev_abc, double* dev_x1x2,
9                          bool* dev_error)
10 {
11     double root = (dev_abc[1] * dev_abc[1]) - (4 * dev_abc[0]
12         * dev_abc[2]);
13     // printf("root: %lf\n", root);
14     if (root < 0) {
15         *dev_error = true;
16     }
17     else {
18         *dev_error = false;
```

```

17     dev_x1x2[0] = ((-1 * dev_abc[1] - sqrt(root)) / (2 *
18         dev_abc[0]));
19     dev_x1x2[1] = ((-1 * dev_abc[1] + sqrt(root)) / (2 *
20         dev_abc[0]));
21 }
22
23 int main() {
24     double* n_host = (double*)malloc(sizeof(double) * 3); //
        not cast, error
25     double* x1x2_host = (double*)malloc(sizeof(double) * 2);
26     bool* error_host = (bool*)malloc(sizeof(bool));
27
28     double* n_dev;
29     double* x1x2_dev;
30     bool* error_dev;
31     cudaMalloc((void**)&n_dev, sizeof(double) * 3);
32     cudaMalloc((void**)&x1x2_dev, sizeof(double) * 2);
33     cudaMalloc((void**)&error_dev, sizeof(bool)); // &bool
        error
34
35     for (int i = 0; i < 3; i++) {
36         printf("%c: ", char(i + 97)); //printf("%s", (i + 65))
            ; exception
37         scanf("%lf", &n_host[i]); // "A:%lf" not error, but
            input incomplete // \n weird results
38     }
39
40     x1x2_host[0] = 0;
41     x1x2_host[1] = 0;
42     *error_host = false;
43
44     cudaMemcpy(n_dev, n_host, sizeof(double) * 3,
        cudaMemcpyHostToDevice);
45     cudaMemcpy(x1x2_dev, x1x2_host, sizeof(double) * 2,
        cudaMemcpyHostToDevice); // not necessary
46     cudaMemcpy(error_dev, error_host, sizeof(bool),
        cudaMemcpyHostToDevice); // not necessary
47
48     solveGPU << < 1, 1 >> > (n_dev, x1x2_dev, error_dev);
49
50     // cout << "cuda ptr " << *error_dev << endl; // no error,
        but execption at runtime
51     cudaMemcpy(error_host, error_dev, sizeof(bool),
        cudaMemcpyDeviceToHost);
52     cudaMemcpy(x1x2_host, x1x2_dev, sizeof(double) * 2,
        cudaMemcpyDeviceToHost);
53     if (*error_host) {
54         printf("GPU Result:\n");
55         printf("The solution does not exist\n");

```

```
56     }
57     else {
58         printf("GPU Result:\n");
59         printf("x1 = %lf x2 = %lf\n", x1x2_host[0], x1x2_host
           [1]);
60     }
61
62 }
```

Other Findings

```
1  #include "cuda_runtime.h"
2  #include "device_launch_parameters.h"
3
4  #include <stdio.h>
5  #include <stdlib.h>
6  #include <math.h>
7  #include <iostream>
8
9  using namespace std;
10
11 __global__ void solveGPU(double* dev_abc, double* dev_x1x2,
    int* dev_error)
12 {
13     double root = (dev_abc[1] * dev_abc[1]) - (4 * dev_abc[0]
        * dev_abc[2]);
14     if (root < 0) {
15         *dev_error = true;
16     }
17     else {
18         *dev_error = false;
19         dev_x1x2[0] = ((-1 * dev_abc[1] - sqrt(root)) / (2 *
            dev_abc[0]));
20         dev_x1x2[1] = ((-1 * dev_abc[1] + sqrt(root)) / (2 *
            dev_abc[0]));
21     }
22
23 }
24
25 int main() {
26     double n_host[3] = { 0 };
27     double x1x2_host[2] = { 0 };
28     bool error_host = false;
29
30     double* n_dev;
31     double* x1x2_dev;
32     int* error_dev; // gives no error
33     cudaMalloc((void**)&n_dev, sizeof(double) * 3);
34     cudaMalloc((void**)&x1x2_dev, sizeof(double) * 2);
```

```

35     cudaMalloc((void**)&error_dev, sizeof(bool));
36
37     for (int i = 0; i < 3; i++) {
38         printf("%c: ", char(i + 97));
39         scanf("%lf", &n_host[i]);
40     }
41
42     cudaMemcpy(n_dev, n_host, sizeof(double) * 3,
43               cudaMemcpyHostToDevice);
44     cudaMemcpy(x1x2_dev, x1x2_host, sizeof(double) * 2,
45               cudaMemcpyHostToDevice); // not necessary
46     cudaMemcpy(error_dev, &error_host, sizeof(bool),
47               cudaMemcpyHostToDevice); // not necessary
48
49     solveGPU << < 1, 1 >> > (n_dev, x1x2_dev, error_dev);
50
51     cudaMemcpy(&error_host, error_dev, sizeof(bool),
52               cudaMemcpyDeviceToHost);
53     cudaMemcpy(x1x2_host, x1x2_dev, sizeof(double) * 2,
54               cudaMemcpyDeviceToHost);
55     if (error_host) {
56         printf("GPU Result:\n");
57         printf("The solution does not exist\n");
58     }
59     else {
60         printf("GPU Result:\n");
61         printf("x1 = %lf x2 = %lf\n", x1x2_host[0], x1x2_host
62               [1]);
63     }
64
65     //free(n_host); // exc
66     //free(x1x2_host); // exc
67     //free(&error_host); // exc
68
69     cudaFree(n_dev);
70     cudaFree(x1x2_dev);
71     cudaFree(error_dev);
72 }

```

```

1  int* test;
2  cudaMalloc((void**)&test, sizeof(bool)); // no error

```

```

1  #include "cuda_runtime.h"
2  #include "device_launch_parameters.h"
3
4  #include <stdio.h>
5  #include <stdlib.h>
6  #include <math.h>
7  #include <iostream>
8

```

```
9  using namespace std;
10
11  __global__ void solveGPU(double* dev_abc, double* dev_x1x2,
12                          int* dev_error)
13  {
14      double root = (dev_abc[1] * dev_abc[1]) - (4 * dev_abc[0]
15          * dev_abc[2]);
16      if (root < 0) {
17          *dev_error = true;
18      }
19      else {
20          *dev_error = false;
21          dev_x1x2[0] = ((-1 * dev_abc[1] - sqrt(root)) / (2 *
22              dev_abc[0]));
23          dev_x1x2[1] = ((-1 * dev_abc[1] + sqrt(root)) / (2 *
24              dev_abc[0]));
25      }
26  }
27
28  int main() {
29      double n_host[3] = { 0 };
30      double x1x2_host[2] = { 0 };
31      bool error_host = false;
32
33      double* n_dev;
34      double* x1x2_dev;
35      int* error_dev; // gives no error
36      cudaMalloc((void**)&n_dev, sizeof(double) * 3);
37      cudaMalloc((void**)&x1x2_dev, sizeof(double) * 2);
38      cudaMalloc((void**)&error_dev, sizeof(bool));
39
40      for (int i = 0; i < 3; i++) {
41          printf("%c: ", char(i + 97));
42          scanf("%lf", &n_host[i]);
43      }
44
45      cudaMemcpy(n_dev, n_host, sizeof(double) * 3,
46          cudaMemcpyHostToDevice);
47      cudaMemcpy(x1x2_dev, x1x2_host, sizeof(double) * 2,
48          cudaMemcpyHostToDevice); // not necessary
49      cudaMemcpy(error_dev, &error_host, sizeof(bool),
50          cudaMemcpyHostToDevice); // not necessary
51
52      solveGPU << < 1, 1 >> > (n_dev, x1x2_dev, error_dev);
53
54      cudaMemcpy(&error_host, error_dev, sizeof(bool),
55          cudaMemcpyDeviceToHost);
56      cudaMemcpy(x1x2_host, x1x2_dev, sizeof(double) * 2,
57          cudaMemcpyDeviceToHost);
```

```
51     if (error_host) {
52         printf("GPU Result:\n");
53         printf("The solution does not exist\n");
54     }
55     else {
56         printf("GPU Result:\n");
57         printf("x1 = %lf x2 = %lf\n", x1x2_host[0], x1x2_host
58             [1]);
59     }
```

- `cudaMalloc(void** devPtr, size_t size)`: Allocates `size` bytes of linear memory on the device and returns in `*devPtr` a pointer to the allocated memory. Memory not cleared.
- `cudaMemcpy (void* dst, const void* src, size_t count, cudaMemcpyKind kind)`: Copies `count` bytes from the memory area pointed to by `src` to the memory area pointed to by `dst`. Calling `cudaMemcpy()` with `dst` and `src` pointers that do not match the direction of the copy results in an undefined behavior.