

RGB Image Manipulation Using CUDA: Practice

To process an image in RGB, we just need to apply what we did for a grayscale image but three times, each for an RGB channel: have a matrix (vector) of data for channel R, G and B.

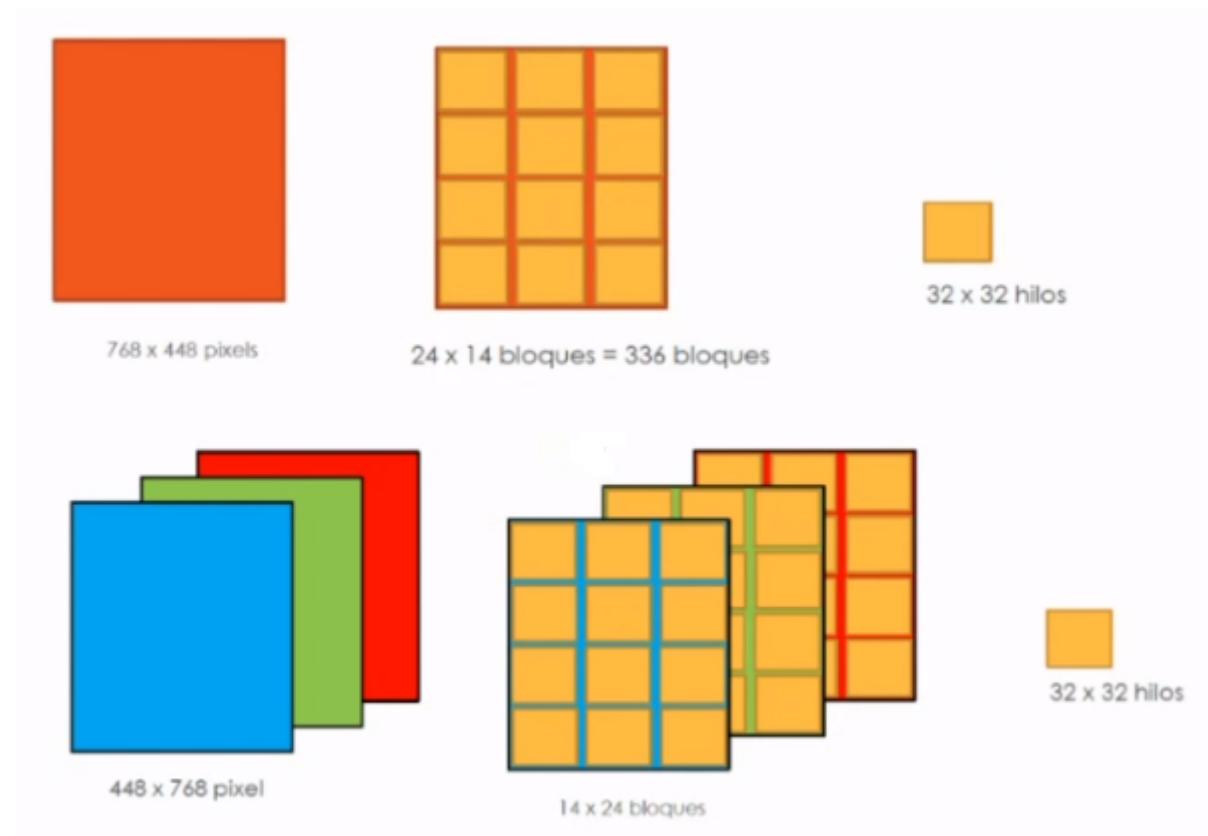


Figure 1: image

Lab 11

Write a program in c/c++ using CUDA in which you implement a kernel to modify the contrast of an RGB image, and another to modify the brightness, considering the following requirements:

- Blocks of 32 x 32 threads
- The kernel for complement, contrast and brightness should be:
 - `__global__ void complement(uchar* R, uchar* G, uchar* B)`
 - `__global__ void contrast(uchar* R, uchar* G, uchar* B, float fc)`
 - `__global__ void brightness(uchar* R, uchar* G, uchar* B, float fb)`
- Include error management with a function:

```
- __host__ void checkCUDAError(const char* msg)
```

Input



Figure 2: img

Solution

```
1  #include "cuda_runtime.h"
2  #include "device_launch_parameters.h"
3
4  #include <stdio.h>
5  #include <stdlib.h>
6  #include <opencv2/opencv.hpp>
7
8  __host__ void checkCUDAError(const char* msg) {
9      cudaError_t error;
10     cudaDeviceSynchronize();
11     error = cudaGetLastError();
12     if (error != cudaSuccess) {
13         printf("ERROR %d: %s (%s)\n", error,
14             cudaGetErrorString(error), msg);
15     }
16 }
17
18 __global__ void complement(uchar* R, uchar* G, uchar* B) {
19     // locate my current block row
20     int threads_per_block = blockDim.x * blockDim.y;
```

```

20     int threads_per_row = threads_per_block * gridDim.x;
21     int row_offset = threads_per_row * blockIdx.y;
22
23     // locate my current block column
24     int block_offset = blockIdx.x * threads_per_block;
25     int threadIdx_inside = blockDim.x * threadIdx.y + threadIdx
        .x;
26
27     int gId = row_offset + block_offset + threadIdx_inside;
28     R[gId] = 255 - R[gId];
29     G[gId] = 255 - G[gId];
30     B[gId] = 255 - B[gId];
31 }
32
33 __global__ void contrast(uchar* R, uchar* G, uchar* B, float
    fc) {
34     // locate my current block row
35     int threads_per_block = blockDim.x * blockDim.y;
36     int threads_per_row = threads_per_block * gridDim.x;
37     int row_offset = threads_per_row * blockIdx.y;
38
39     // locate my current block column
40     int block_offset = blockIdx.x * threads_per_block;
41     int threadIdx_inside = blockDim.x * threadIdx.y + threadIdx
        .x;
42
43     int gId = row_offset + block_offset + threadIdx_inside;
44     if (fc * R[gId] > 255) { R[gId] = 255; } else { R[gId] =
        fc * R[gId]; }
45     if (fc * G[gId] > 255) { G[gId] = 255; } else { G[gId] =
        fc * G[gId]; }
46     if (fc * B[gId] > 255) { B[gId] = 255; } else { B[gId] =
        fc * B[gId]; }
47 }
48
49 __global__ void brightness(uchar* R, uchar* G, uchar* B, float
    fb) {
50     // locate my current block row
51     int threads_per_block = blockDim.x * blockDim.y;
52     int threads_per_row = threads_per_block * gridDim.x;
53     int row_offset = threads_per_row * blockIdx.y;
54
55     // locate my current block column
56     int block_offset = blockIdx.x * threads_per_block;
57     int threadIdx_inside = blockDim.x * threadIdx.y + threadIdx
        .x;
58
59     int gId = row_offset + block_offset + threadIdx_inside;
60     if (fb >= 0){
61     if (fb + R[gId] > 255) { R[gId] = 255; } else { R[gId] =
        fb + R[gId]; }

```

```
62     if (fb + G[gId] > 255) { G[gId] = 255; } else { G[gId] =
        fb + G[gId]; }
63     if (fb + B[gId] > 255) { B[gId] = 255; } else { B[gId] =
        fb + B[gId]; }
64 }
65
66     if (fb < 0) {
67         if (fb + R[gId] < 0) { R[gId] = 0; } else { R[gId] = fb +
            R[gId]; }
68         if (fb + G[gId] < 0) { G[gId] = 0; } else { G[gId] = fb +
            G[gId]; }
69         if (fb + B[gId] < 0) { B[gId] = 0; } else { B[gId] = fb +
            B[gId]; }
70     }
71 }
72
73 using namespace cv;
74 int main() {
75
76     Mat img = imread("antenaRGB.jpg");
77
78     const int R = img.rows;
79     const int C = img.cols;
80
81     Mat imgComp(img.rows, img.cols, img.type());
82     Mat imgCont(img.rows, img.cols, img.type());
83     Mat imgBright(img.rows, img.cols, img.type());
84     uchar* host_r, * host_g, * host_b, * dev_r1, * dev_g1, *
        dev_b1, * dev_r2, * dev_g2, * dev_b2, * dev_r3, *
        dev_g3, * dev_b3;
85     host_r = (uchar*)malloc(sizeof(uchar) * R * C);
86     host_g = (uchar*)malloc(sizeof(uchar) * R * C);
87     host_b = (uchar*)malloc(sizeof(uchar) * R * C);
88
89     cudaMalloc((void**)&dev_r1, sizeof(uchar) * R * C);
90     checkCUDAError("Error at malloc dev_r1");
91     cudaMalloc((void**)&dev_g1, sizeof(uchar) * R * C);
92     checkCUDAError("Error at malloc dev_g1");
93     cudaMalloc((void**)&dev_b1, sizeof(uchar) * R * C);
94     checkCUDAError("Error at malloc dev_b1");
95
96     cudaMalloc((void**)&dev_r2, sizeof(uchar) * R * C);
97     checkCUDAError("Error at malloc dev_r2");
98     cudaMalloc((void**)&dev_g2, sizeof(uchar) * R * C);
99     checkCUDAError("Error at malloc dev_g2");
100    cudaMalloc((void**)&dev_b2, sizeof(uchar) * R * C);
101    checkCUDAError("Error at malloc dev_b2");
102
103    cudaMalloc((void**)&dev_r3, sizeof(uchar) * R * C);
104    checkCUDAError("Error at malloc dev_r3");
105    cudaMalloc((void**)&dev_g3, sizeof(uchar) * R * C);
```

```

106     checkCUDAError("Error at malloc dev_g3");
107     cudaMalloc((void**)&dev_b3, sizeof(uchar) * R * C);
108     checkCUDAError("Error at malloc dev_b3");
109
110     // matrix as vector
111     for (int i = 0; i < R; i++) {
112         for (int j = 0; j < C; j++) {
113             Vec3b pix = img.at<Vec3b>(i, j);
114             host_r[i * C + j] = pix[2];
115             host_g[i * C + j] = pix[1];
116             host_b[i * C + j] = pix[0];
117         }
118     }
119     cudaMemcpy(dev_r1, host_r, sizeof(uchar) * R * C,
120               cudaMemcpyHostToDevice);
121     checkCUDAError("Error at memcpy host_r -> dev_r1");
122     cudaMemcpy(dev_g1, host_g, sizeof(uchar) * R * C,
123               cudaMemcpyHostToDevice);
124     checkCUDAError("Error at memcpy host_r -> dev_g1");
125     cudaMemcpy(dev_b1, host_b, sizeof(uchar) * R * C,
126               cudaMemcpyHostToDevice);
127     checkCUDAError("Error at memcpy host_r -> dev_b1");
128
129     cudaMemcpy(dev_r2, host_r, sizeof(uchar) * R * C,
130               cudaMemcpyHostToDevice);
131     checkCUDAError("Error at memcpy host_r -> dev_r2");
132     cudaMemcpy(dev_g2, host_g, sizeof(uchar) * R * C,
133               cudaMemcpyHostToDevice);
134     checkCUDAError("Error at memcpy host_r -> dev_g2");
135     cudaMemcpy(dev_b2, host_b, sizeof(uchar) * R * C,
136               cudaMemcpyHostToDevice);
137     checkCUDAError("Error at memcpy host_r -> dev_b2");
138
139     cudaMemcpy(dev_r3, host_r, sizeof(uchar) * R * C,
140               cudaMemcpyHostToDevice);
141     checkCUDAError("Error at memcpy host_r -> dev_r3");
142     cudaMemcpy(dev_g3, host_g, sizeof(uchar) * R * C,
143               cudaMemcpyHostToDevice);
144     checkCUDAError("Error at memcpy host_r -> dev_g3");
145     cudaMemcpy(dev_b3, host_b, sizeof(uchar) * R * C,
146               cudaMemcpyHostToDevice);
147     checkCUDAError("Error at memcpy host_r -> dev_b3");
148
149     dim3 block(32, 32);
150     dim3 grid(C / 32, R / 32);
151
152     complement << < grid, block >> > (dev_r1, dev_g1, dev_b1);
153     cudaDeviceSynchronize();
154     checkCUDAError("Error at kernel complement");
155

```

```

147     cudaMemcpy(host_r, dev_r1, sizeof(uchar) * R * C,
148               cudaMemcpyDeviceToHost);
149     cudaMemcpy(host_g, dev_g1, sizeof(uchar) * R * C,
150               cudaMemcpyDeviceToHost);
151     cudaMemcpy(host_b, dev_b1, sizeof(uchar) * R * C,
152               cudaMemcpyDeviceToHost);
153     checkCUDAError("Error at memcpy host_r <- dev_b1");
154     for (int i = 0; i < R; i++) {
155         for (int j = 0; j < C; j++) {
156             imgComp.at<Vec3b>(i, j)[0] = host_b[i * C + j];
157             imgComp.at<Vec3b>(i, j)[1] = host_g[i * C + j];
158             imgComp.at<Vec3b>(i, j)[2] = host_r[i * C + j];
159         }
160     }
161
162     contrast << < grid, block >> > (dev_r2, dev_g2, dev_b2,
163                                     0.5);
164     cudaDeviceSynchronize();
165     checkCUDAError("Error at kernel contrast");
166
167     cudaMemcpy(host_r, dev_r2, sizeof(uchar) * R * C,
168               cudaMemcpyDeviceToHost);
169     cudaMemcpy(host_g, dev_g2, sizeof(uchar) * R * C,
170               cudaMemcpyDeviceToHost);
171     cudaMemcpy(host_b, dev_b2, sizeof(uchar) * R * C,
172               cudaMemcpyDeviceToHost);
173     checkCUDAError("Error at memcpy host_r <- dev_b2");
174
175     for (int i = 0; i < R; i++) {
176         for (int j = 0; j < C; j++) {
177             imgCont.at<Vec3b>(i, j)[0] = host_b[i * C + j];
178             imgCont.at<Vec3b>(i, j)[1] = host_g[i * C + j];
179             imgCont.at<Vec3b>(i, j)[2] = host_r[i * C + j];
180         }
181     }
182
183     brightness << < grid, block >> > (dev_r3, dev_g3, dev_b3,
184                                       100);
185     cudaDeviceSynchronize();
186     checkCUDAError("Error at kernel brightness");
187
188     cudaMemcpy(host_r, dev_r3, sizeof(uchar) * R * C,
189               cudaMemcpyDeviceToHost);
190     cudaMemcpy(host_g, dev_g3, sizeof(uchar) * R * C,
191               cudaMemcpyDeviceToHost);
192     cudaMemcpy(host_b, dev_b3, sizeof(uchar) * R * C,
193               cudaMemcpyDeviceToHost);

```

```
188     checkCUDAError("Error at memcpy host_r <- dev_g3");
189     cudaMemcpy(host_b, dev_b3, sizeof(uchar) * R * C,
190               cudaMemcpyDeviceToHost);
191     checkCUDAError("Error at memcpy host_r <- dev_b3");
192
193     for (int i = 0; i < R; i++) {
194         for (int j = 0; j < C; j++) {
195             imgBright.at<Vec3b>(i, j)[0] = host_b[i * C + j];
196             imgBright.at<Vec3b>(i, j)[1] = host_g[i * C + j];
197             imgBright.at<Vec3b>(i, j)[2] = host_r[i * C + j];
198         }
199     }
200
201     imshow("Image", img);
202     imshow("Image Complement", imgComp);
203     imshow("Image Contrast", imgCont);
204     imshow("Image Brightness", imgBright);
205     waitKey(0);
206
207     free(host_r);
208     free(host_g);
209     free(host_b);
210     cudaFree(dev_r1);
211     cudaFree(dev_g1);
212     cudaFree(dev_b1);
213     cudaFree(dev_r2);
214     cudaFree(dev_g2);
215     cudaFree(dev_b2);
216     cudaFree(dev_r3);
217     cudaFree(dev_g3);
218     cudaFree(dev_b3);
219
220     return 0;
221 }
```


Output

Complement



Figure 3: img


Contrast

Figure 4: img

Brightness



Figure 5: img