

## Kernel Validation

Verify just once that the kernel operations are correctly performed, which is crucial when working with large amounts of data. This consists in a comparison between the result of the same operation done both in CPU and GPU.

We would need to add a function that does that same pixel operation in the host, to continue the same examples with image processing, for example. This is basically making the kernel function but in the host.

Validation: execute the same kernel operations in the host and compare the results given by the CPU and GPU.

## Exercise

Execute a kernel validation for the last kernel where we made the complement of an RGB image. The kernel validation includes two functions: one for performing the complement in the CPU and one for comparing the CPU image and the GPU image.

## Solution

```
1  #include "cuda_runtime.h"
2  #include "device_launch_parameters.h"
3
4  #include <stdio.h>
5  #include <stdlib.h>
6  #include <opencv2/opencv.hpp>
7
8  using namespace cv;
9
10 __host__ void checkCUDAError(const char* msg) {
11     cudaError_t error;
12     cudaDeviceSynchronize();
13     error = cudaGetLastError();
14     if (error != cudaSuccess) {
15         printf("ERROR %d: %s (%s)\n", error,
16             cudaGetErrorString(error), msg);
17     }
18 }
19
20 __global__ void complement(uchar* RGB) {
21     // locate my current block row
22     int threads_per_block = blockDim.x * blockDim.y;
23     int threads_per_row = threads_per_block * gridDim.x;
24     int row_offset = threads_per_row * blockIdx.y;
```

```
25
26     // locate my current block column
27     int block_offset = blockIdx.x * threads_per_block;
28     int threadIdx_inside = blockDim.x * threadIdx.y + threadIdx
        .x;
29
30     // locate my current grid row
31     int thread_per_grid = (gridDim.x * gridDim.y *
        threads_per_block);
32     int gridOffset = blockIdx.z * thread_per_grid;
33
34     int gId = gridOffset + row_offset + block_offset +
        threadIdx_inside;
35     RGB[gId] = 255 - RGB[gId];
36 }
37
38 __host__ void complementCPU(Mat* original, Mat* comp) {
39     for (int i = 0; i < original->rows; i++) {
40         for (int j = 0; j < original->cols; j++) {
41             comp->at<Vec3b>(i, j)[0] = 255 - original->at<
                Vec3b>(i, j)[0];
42             comp->at<Vec3b>(i, j)[1] = 255 - original->at<
                Vec3b>(i, j)[1];
43             comp->at<Vec3b>(i, j)[2] = 255 - original->at<
                Vec3b>(i, j)[2];
44         }
45     }
46 }
47
48 __host__ bool validationKernel(Mat img1, Mat img2) {
49     Vec3b* pImg1, * pImg2;
50     for (int k = 0; k < 3; k++) {
51         for (int i = 0; i < img1.rows; i++) {
52             pImg1 = img1.ptr<Vec3b>(i);
53             pImg2 = img2.ptr<Vec3b>(i);
54             for (int j = 0; j < img1.cols; j++) {
55                 if (pImg1[j][k] != pImg2[j][k]) {
56                     printf("Error at kernel validation\n");
57                     return true;
58                 }
59             }
60         }
61     }
62     printf("Kernel validation successful\n");
63     return false;
64 }
65
66 int main() {
67
68     Mat img = imread("antenaRGB.jpg");
69
```

```

70     const int R = img.rows;
71     const int C = img.cols;
72
73     Mat imgComp(img.rows, img.cols, img.type());
74     Mat imgCompCPU(img.rows, img.cols, img.type());
75     uchar* host_rgb, * dev_rgb;
76     host_rgb = (uchar*)malloc(sizeof(uchar) * R * C * 3);
77
78     cudaMalloc((void**)&dev_rgb, sizeof(uchar) * R * C * 3);
79     checkCUDAError("Error at malloc dev_r1");
80
81     // matrix as vector
82     for (int k = 0; k < 3; k++) {
83         for (int i = 0; i < R; i++) {
84             for (int j = 0; j < C; j++) {
85                 Vec3b pix = img.at<Vec3b>(i, j);
86
87                 host_rgb[i * C + j + (k * R * C)] = pix[k];
88
89             }
90         }
91     }
92     cudaMemcpy(dev_rgb, host_rgb, sizeof(uchar) * R * C * 3,
93               cudaMemcpyHostToDevice);
94     checkCUDAError("Error at memcpy host_rgb -> dev_rgb");
95
96     dim3 block(32, 32);
97     dim3 grid(C / 32, R / 32, 3);
98
99     complement << < grid, block >> > (dev_rgb);
100    cudaDeviceSynchronize();
101    checkCUDAError("Error at kernel complement");
102
103    cudaMemcpy(host_rgb, dev_rgb, sizeof(uchar) * R * C * 3,
104              cudaMemcpyDeviceToHost);
105    checkCUDAError("Error at memcpy host_rgb <- dev_rgb");
106
107    for (int k = 0; k < 3; k++) {
108        for (int i = 0; i < R; i++) {
109            for (int j = 0; j < C; j++) {
110                imgComp.at<Vec3b>(i, j)[k] = host_rgb[i * C +
111                j + (k * R * C)];
112            }
113        }
114    }
115
116    complementCPU(&img, &imgCompCPU);
117    bool error = validationKernel(imgCompCPU, imgComp);
118
119    if (error) {
120        printf("Check kernel operations\n");

```

```
118         return 0;
119     }
120
121
122     imshow("Image", img);
123     imshow("Image Complement CPU", imgCompCPU);
124     imshow("Image Complement GPU", imgComp);
125     waitKey(0);
126
127     free(host_rgb);
128     cudaFree(dev_rgb);
129
130     return 0;
131 }
```