

Practice

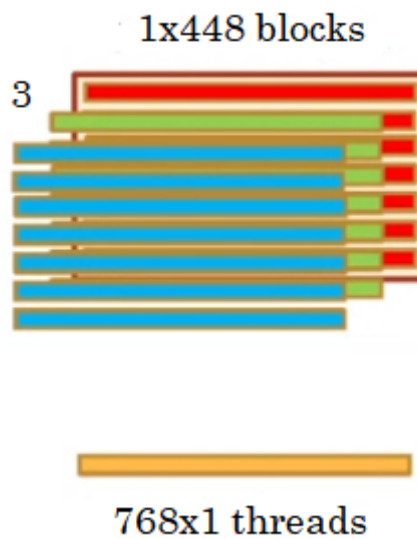


Figure 1: img

Lab 13

Write a program in c/c++ using CUDA in which you implement a kernel to calculate an RGB image complement. The kernel must be verified, consider the requirements:

- The complement of an image is defined as:

$$I(x,y) = 255 - I(x,y)$$

- Blocks of 768 x 1 threads.
- A grid of 1 x 448 x 3 blocks.
- The kernel signature should be: `__global__ void complement(uchar* RGB)`
- The CPU complement function signature should be: `__host__ void complementCPU(Mat* original, Mat* comp)`
- The kernel validation signature should be: `__host__ bool validationKernel(Mat img1, Mat img2)`

Solution

```
1 #include "cuda_runtime.h"
2 #include "device_launch_parameters.h"
3
4 #include <stdio.h>
```

```
5  #include <stdlib.h>
6  #include <opencv2/opencv.hpp>
7
8  using namespace cv;
9
10 __host__ void checkCUDAError(const char* msg) {
11     cudaError_t error;
12     cudaDeviceSynchronize();
13     error = cudaGetLastError();
14     if (error != cudaSuccess) {
15         printf("ERROR %d: %s (%s)\n", error,
16             cudaGetErrorString(error), msg);
17     }
18 }
19
20 __global__ void complement(uchar* RGB) {
21     // locate my current block row
22     int threads_per_block = blockDim.x;
23     int threads_per_row = threads_per_block * gridDim.x;
24     int row_offset = threads_per_row * blockIdx.y;
25
26     // locate my current block column
27     int block_offset = blockIdx.x * threads_per_block;
28
29     // locate my current grid row
30     int thread_per_grid = (gridDim.x * gridDim.y *
31         threads_per_block);
32     int gridOffset = blockIdx.z * thread_per_grid;
33
34     int gId = gridOffset + row_offset + block_offset +
35         threadIdx.x;
36     RGB[gId] = 255 - RGB[gId];
37 }
38
39 __host__ void complementCPU(Mat* original, Mat* comp) {
40     for (int i = 0; i < original->rows; i++) {
41         for (int j = 0; j < original->cols; j++) {
42             comp->at<Vec3b>(i, j)[0] = 255 - original->at<
43                 Vec3b>(i, j)[0];
44             comp->at<Vec3b>(i, j)[1] = 255 - original->at<
45                 Vec3b>(i, j)[1];
46             comp->at<Vec3b>(i, j)[2] = 255 - original->at<
47                 Vec3b>(i, j)[2];
48         }
49     }
50 }
51
52 __host__ bool validationKernel(Mat img1, Mat img2) {
53     Vec3b* pImg1, * pImg2;
54     for (int k = 0; k < 3; k++) {
```

```

50     for (int i = 0; i < img1.rows; i++) {
51         pImg1 = img1.ptr<Vec3b>(i);
52         pImg2 = img2.ptr<Vec3b>(i);
53         for (int j = 0; j < img1.cols; j++) {
54             if (pImg1[j][k] != pImg2[j][k]) {
55                 printf("Error at kernel validation\n");
56                 return true;
57             }
58         }
59     }
60 }
61 printf("Kernel validation successful\n");
62 return false;
63 }
64
65 int main() {
66
67     Mat img = imread("antenaRGB.jpg");
68
69     const int R = img.rows;
70     const int C = img.cols;
71
72     Mat imgComp(img.rows, img.cols, img.type());
73     Mat imgCompCPU(img.rows, img.cols, img.type());
74     uchar* host_rgb, * dev_rgb;
75     host_rgb = (uchar*)malloc(sizeof(uchar) * R * C * 3);
76
77     cudaMalloc((void**)&dev_rgb, sizeof(uchar) * R * C * 3);
78     checkCUDAError("Error at malloc dev_r1");
79
80     // matrix as vector
81     for (int k = 0; k < 3; k++) {
82         for (int i = 0; i < R; i++) {
83             for (int j = 0; j < C; j++) {
84                 Vec3b pix = img.at<Vec3b>(i, j);
85
86                 host_rgb[i * C + j + (k * R * C)] = pix[k];
87
88             }
89         }
90     }
91     cudaMemcpy(dev_rgb, host_rgb, sizeof(uchar) * R * C * 3,
92               cudaMemcpyHostToDevice);
93     checkCUDAError("Error at memcpy host_rgb -> dev_rgb");
94
95     //dim3 block(32, 32);
96     //dim3 grid(C / 32, R / 32, 3); // 24 14
97     dim3 block(C, 1, 1); // 768
98     dim3 grid(1, R, 3); // 448
99
100    complement << < grid, block >> > (dev_rgb);

```

```
100     cudaDeviceSynchronize();
101     checkCUDAError("Error at kernel complement");
102
103     cudaMemcpy(host_rgb, dev_rgb, sizeof(uchar) * R * C * 3,
104               cudaMemcpyDeviceToHost);
105     checkCUDAError("Error at memcpy host_rgb <- dev_rgb");
106
107     for (int k = 0; k < 3; k++) {
108         for (int i = 0; i < R; i++) {
109             for (int j = 0; j < C; j++) {
110                 imgComp.at<Vec3b>(i, j)[k] = host_rgb[i * C +
111                 j + (k * R * C)];
112             }
113         }
114     }
115
116     complementCPU(&img, &imgCompCPU);
117     bool error = validationKernel(imgCompCPU, imgComp);
118
119     if (error) {
120         printf("Check kernel operations\n");
121         return 0;
122     }
123
124     imshow("Image", img);
125     imshow("Image Complement CPU", imgCompCPU);
126     imshow("Image Complement GPU", imgComp);
127     waitKey(0);
128
129     free(host_rgb);
130     cudaFree(dev_rgb);
131
132     return 0;
133 }
```

Input




Figure 2: img

Output




Figure 3: img




Figure 4: img